




CS5491: Artificial Intelligence

Reinforcement Learning



Instructor: Kai Wang

Today



Why is
Reinforcement
Learning?

Markov Decision
Processes


Algorithm
Anatomy

Policy Gradients

Today



Why is
Reinforcement
Learning?



Markov Decision
Processes




Algorithm
Anatomy



Policy Gradients

Reinforcement Learning





Example: Driving a Car

- ❖ Objective: drive a car not to hit the lion
- ❖ State: road scene, status of lions (why not percepts?)
- ❖ Action: turn left or right
- ❖ Reward: 1 at each time step if the lion is not hit.



Example: Robot Locomotion

- ❖ Objective: make the robot move forward
- ❖ State: angle and position of the joints
- ❖ Action: torques applied on joints
- ❖ Reward: 1 at each time step upright + forward movement





Example: Atari Games

- ❖ Objective: complete the game with the highest score
- ❖ State: raw pixel inputs of the game state
- ❖ Action: game controls (left, right, up, down)
- ❖ Reward: score increase/decrease at each time step



Today




Why is
Reinforcement
Learning?

Markov Decision
Processes

Algorithm
Anatomy

Policy Gradients

Markov Decision Process



Markov Decision Process provides a mathematical formulation of the Reinforcement Learning problem, defined as $(\mathcal{S}, \mathcal{A}, \mathcal{R}, \mathbb{P}, \gamma)$.

- ❖ \mathcal{S} : set of all possible states
- ❖ \mathcal{A} : set of all possible actions
- ❖ \mathcal{R} : distribution of reward given (state, action) pair
- ❖ \mathbb{P} : transition probability, i.e., distribution over next state given (state, action) pair
- ❖ γ : discounting factor

Markov Decision Process

- ❖ Step $t = 0$: environment samples initial state $s_0 \sim p(s_0)$
- ❖ For $t = 0$ until done:
 - Agent selects action a_t based on a policy $\pi(a_t | s_t)$
 - Environment samples reward $r_t \sim R(\cdot | s_t, a_t)$
 - Environment samples next state $s_{t+1} \sim P(\cdot | s_t, a_t)$
 - Agent receives reward r_t and next state s_{t+1}
- ❖ Objective: find policy π^* that maximizes cumulative discounted reward: $\sum_{t \geq 0} \gamma^t r_t$.

The Optimal Policy

- ❖ How do we handle the randomness (initial state, transition probability ...)?
 - Maximize the ~~expected sum of rewards!~~
 - $\pi^* = \operatorname{argmax}_\pi \mathbb{E}[\sum_{t \geq 0} \gamma^t r_t | \pi]$ with
 $s_0 \sim p(s_0), a_t \sim \pi(\cdot | s_t), s_{t+1} \sim p(\cdot | s_t, a_t)$
- ❖ Why discounting factor?

The Optimal Policy




We want?

Goin...eyon...

- Bandits
 - No state
 - previous
 - irrelevant
 - simple
- Finite Markets
 - Finite
 - Agents
 - “Actions”
 - instances


Value Function



Following a policy produces sample trajectories $(s_0, a_0, r_0, s_1, a_1, r_1, \dots)$. The value function evaluates how good a state s is by measuring the expected cumulative reward from following the policy from state s .

$$V^\pi(s) = \mathbb{E}\left[\sum_{t \geq 0} \gamma^t r_t | s_0 = s, \pi\right]$$

Q-value Function



Following a policy produces sample trajectories $(s_0, a_0, r_0, s_1, a_1, r_1, \dots)$. The Q-value function evaluates how good a state-action pair (s, a) is by measuring the expected cumulative reward from taking action a in state s and following the policy.

$$Q^\pi(s, a) = \mathbb{E}\left[\sum_{t \geq 0} \gamma^t r_t | s_0 = s, a_0 = a, \pi\right]$$

Value Function and Q-value Function

◆ Connections

$$V^\pi(s) = E_{a \sim \pi(a|s)} Q^\pi(s, a)$$

◆ Using Q-functions and value functions

- If we have policy π , and we know $Q^\pi(s, a)$, then we can improve π :

Set $\pi'(a|s) = 1$, if $a = \operatorname{argmax}_a Q^\pi(s, a)$


This policy is at least as good as π (and probably better)!

- Compute gradient to increase the probability of good actions a :
If $Q^\pi(s, a) > V^\pi(s)$, then a is better than average!


Today



Why is
Reinforcement
Learning?



Markov Decision
Processes




Algorithm
Anatomy



Policy Gradients


Types of RL Algorithms

❖ Value function based algorithms



Types of RL Algorithms

❖ Direct policy gradients algorithms




evaluate returns

$$R_\tau = \sum_t r(s_t, a_t)$$

$$\theta \leftarrow \theta + \alpha \sum_{t \geq 0} r(\tau) \nabla_\theta \log \pi_\theta(a_t | s_t)$$


Types of RL Algorithms

❖ Actor-critic algorithms



Types of RL Algorithms

❖ Model-based algorithms



Model-based Algorithms

❖ A few options


- Just use the model to plan (no policy)
Discrete planning in discrete action spaces, e.g., Monte Carlo tree search
- Backpropagation gradients into the policy
Requires some tricks to make it work
- Use the model to learn a value function
Dynamic programming

Tradeoffs Between Algorithms


- ❖ Different Tradeoffs
 - Sample efficiency
 - Stability & easy of use
- ❖ Different assumptions
 - Stochastic or deterministic?
 - Continuous or discrete?
 - Episodic or infinite horizon?
- ❖ Different things are easy or hard in different settings
 - Easier to represent the policy?
 - Easier to represent the model?

Comparison on Sample Efficiency

- ❖ Sample efficiency = how many samples do we need to get a good policy?
- ❖ Most important question: is the algorithm off-policy?
 - Off-policy: able to improve the policy without generating new samples from that policy
 - On policy: each time the policy is changed, even a little bit, we need to generate new samples



Comparison on Sample Efficiency



Why should we use a less efficient algorithm?

Wall clock time is not the same as sample efficiency!

Comparison on Stability and Ease of Use

- ❖ Does it converge?
- ❖ And if it converges, to what?
- ❖ And does it converge every time?

Why is any of this even a question???

- ❖ Supervised learning: almost always gradient descent
- ❖ Reinforcement learning: often not gradient descent
 - Q-learning: fixed point iteration
 - Model-based RL: model is not optimized for expected reward
 - Policy gradient: is gradient descent, but also often the least efficient!

Comparison on Stability and Ease of Use



- ❖ Value function fitting
 - At best, minimizes error of fit (“Bellman error”)
Not the same as expected reward
 - At worst, doesn’t optimize anything
Many popular deep RL value fitting algorithms are not guaranteed to converge to anything in the non-linear case
- ❖ Model-based RL
 - Model minimizes error of fit -> this will converge
 - No guarantee that better model = better policy
- ❖ Policy gradient
 - The only one that actually performs gradient descent on the true objective.

Comparison on Assumptions

- ❖ Common assumption #1: full observability
 - Generally assumed by value function fitting methods
 - Can be mitigated by adding recurrence

- ❖ Common assumption #2: episodic learning
 - Often assumed by pure policy gradient methods
 - Assumed by some model-based RL methods

- ❖ Common assumption #3: continuity or smoothness
 - Assumed by some continuous value function methods
 - Often assumed by some model-based RL methods




Examples of Algorithms

- ❖ Value function fitting methods
 - Q-learning, DQN
 - Fitted value iteration
- ❖ Policy gradient methods
 - REINFORCE
 - Trust region policy optimization
- ❖ Actor-critic algorithms
 - Asynchronous advantage actor-critic (A3C)
- ❖ Model-based RL algorithms
 - Dyna


Today



Why is
Reinforcement
Learning?



Markov Decision
Processes



Algorithm
Anatomy



Policy Gradients

Policy Gradients

- ❖ We define a class of parameterized policies $\Pi = \{\pi_\theta, \theta \in \mathbb{R}^m\}$
- ❖ Objective: find the optimal policy $\theta^* = \arg \max_{\theta} J(\theta)$

$$J(\theta) = \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r^t \mid \pi_\theta \right]$$

REINFORCE

- ❖ Mathematically, we can write

$$J(\theta) = \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r^t | \pi_\theta \right] = \mathbb{E}_{\tau \sim p(\tau; \theta)} [r(\tau)]$$

$$= \int_{\tau} r(\tau) p(\tau; \theta) d\tau$$

- ❖ $r(\tau)$ is the reward of a trajectory $\tau = (s_0, a_0, r_0, s_1, \dots)$

REINFORCE

- ❖ Expected Reward

$$J(\theta) = \mathbb{E}_{\tau \sim p(\tau; \theta)}[r(\tau)] = \int_{\tau} r(\tau)p(\tau; \theta)d\tau$$


- ❖ Differentiate this

$$p(\tau; \theta)\nabla_{\theta} \log p(\tau; \theta) = \nabla_{\theta} p(\tau; \theta)$$

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \int_{\tau} r(\tau)\nabla_{\theta} p(\tau; \theta)d\tau \\ &= \int_{\tau} r(\tau)p(\tau; \theta)\nabla_{\theta} \log p(\tau; \theta) d\tau \\ &= \mathbb{E}_{\tau \sim p(\tau; \theta)}[r(\tau)\nabla_{\theta} \log p(\tau; \theta)]\end{aligned}$$

REINFORCE

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau; \theta)} [r(\tau) \nabla_{\theta} \log p(\tau; \theta)]$$



We have $p(\tau; \theta) = \prod_{t \geq 0} p(s_{t+1} | s_t, a_t) \pi_{\theta}(a_t | s_t)$

Thus $\log p(\tau; \theta) = \sum_{t \geq 0} \log p(s_{t+1} | s_t, a_t) + \log \pi_{\theta}(a_t | s_t)$

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i^N \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_{it} | s_{it})$$

Intuition:

If $r(\tau)$ is high, push up the probabilities of the actions seen

If $r(\tau)$ is low, push down the probabilities of the actions seen

REINFORCE

function REINFORCE

 Initialise θ arbitrarily

for each episode $\{s_1, a_1, r_2, \dots, s_{T-1}, a_{T-1}, r_T\} \sim \pi_\theta$ **do**

for $t = 1$ to $T - 1$ **do**

$$\theta \leftarrow \theta + \alpha \nabla_\theta \log \pi_\theta(s_t, a_t) v_t$$

end for

end for

return θ

end function



$$v_t = \sum_t^T r_t$$

Problems with the Policy Gradient

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i^N \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_{it} | s_{it})$$

High variance!

Even worse: what if the two good samples have $r(\tau) = 0$?



Reducing Variance

❖ Using causality

Policy at time t' cannot affect reward at time t when $t < t'$

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_{it} | s_{it})$$



$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i \sum_{t \geq 0} \nabla_{\theta} \log \pi_{\theta}(a_{it} | s_{it}) \sum_{t'=t}^T r(s_{it'}, a_{it'})$$

Baselines

❖ Using a baseline

Subtracting a baseline is unbiased in expectation!


Average reward is not the best baseline, but it's pretty good!

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i^N \sum_{t \geq 0} (r(\tau) - b) \nabla_{\theta} \log \pi_{\theta}(a_{it} | s_{it})$$

$$b = \frac{1}{N} \sum_i^N r(\tau)$$

$$E[\nabla_{\theta} \log p_{\theta}(\tau) b] = \int p_{\theta}(\tau) \nabla_{\theta} \log p_{\theta}(\tau) b d\tau = \int \nabla_{\theta} p_{\theta}(\tau) b d\tau = b \nabla_{\theta} \int p_{\theta}(\tau) d\tau = b \nabla_{\theta} 1 = 0$$

$$p(\tau; \theta) \nabla_{\theta} \log p(\tau; \theta) = \nabla_{\theta} p(\tau; \theta)$$



Optimal Baseline

$$\text{Var}[x] = E[x^2] - E[x]^2$$

$$\nabla_{\theta} J(\theta) = E_{\tau \sim p_{\theta}(\tau)} [\nabla_{\theta} \log p_{\theta}(\tau) (r(\tau) - b)]$$

$$\text{Var} = E_{\tau \sim p_{\theta}(\tau)} [(\nabla_{\theta} \log p_{\theta}(\tau) (r(\tau) - b))^2] - E_{\tau \sim p_{\theta}(\tau)} [\nabla_{\theta} \log p_{\theta}(\tau) (r(\tau) - b)]^2$$

$$\begin{aligned} \frac{d\text{Var}}{db} &= \frac{d}{db} E[g(\tau)^2 (r(\tau) - b)^2] = \frac{d}{db} \left(\cancel{E[g(\tau)^2 r(\tau)^2]} - 2E[g(\tau)^2 r(\tau)b] + b^2 E[g(\tau)^2] \right) \\ &= -2E[g(\tau)^2 r(\tau)] + 2bE[g(\tau)^2] = 0 \end{aligned}$$

$$b = \frac{E[g(\tau)^2 r(\tau)]}{E[g(\tau)^2]}$$

This is just expected reward, but weighted by gradient magnitudes!

Off-policy Policy Gradients

- ❖ On policy learning can be extremely inefficient!

```
function REINFORCE
    Initialise  $\theta$  arbitrarily
    for each episode  $\{s_1, a_1, r_2, \dots, s_{T-1}, a_{T-1}, r_T\} \sim \pi_\theta$  do
        for  $t = 1$  to  $T - 1$  do
             $\theta \leftarrow \theta + \alpha \nabla_\theta \log \pi_\theta(s_t, a_t) v_t$ 
        end for
    end for
    return  $\theta$ 
end function
```

Off-policy Policy Gradients


- ◆ We have off-policy samples $p(\tau; \bar{\theta})$ (i.e., $\bar{p}(\tau)$) instead

$$J(\theta) = E_{\tau \sim \bar{p}(\tau)} \left[\frac{p_\theta(\tau)}{\bar{p}(\tau)} r(\tau) \right]$$

$$p_\theta(\tau) = p(s_1) \prod_{t=1}^T \pi_\theta(a_t | s_t) p(s_{t+1} | s_t, a_t)$$

$$\frac{p_\theta(\tau)}{\bar{p}(\tau)} = \frac{p(s_1) \prod_{t=1}^T \pi_\theta(a_t | s_t) p(s_{t+1} | s_t, a_t)}{p(s_1) \prod_{t=1}^T \bar{\pi}(a_t | s_t) p(s_{t+1} | s_t, a_t)} = \frac{\prod_{t=1}^T \pi_\theta(a_t | s_t)}{\prod_{t=1}^T \bar{\pi}(a_t | s_t)}$$

RL Applications



Goals

-  Understand what Reinforcement Learning (RL) is.
-  Understand categorization of existing Reinforcement Learning algorithms.
-  Learn how to use Keras to implement REINFORCE.
-  Learn how to debug REINFORCE and make it work in practice.

Important This Week

- ✓ Read [Reinforcement Learning: An Introduction.](#)
- ✓ Know more about Reinforcement Learning algorithms here.
<http://rail.eecs.berkeley.edu/deeprlcourse/>
- ✓ Know more about implementation issues of RL here.
<https://spinningup.openai.com/en/latest/>